

HarvardX Data Science Capstone Movie recommendation script

Codrin Kruijne

30/5/2019

```
# PLEASE NOTE THIS SCRIPT HAS ONLY BEEN TESTED ON A 6-CORE, 64GB MEMORY MACHINE AND WILL PROBABLY NOT RUN  
parallel::detectCores()
```

```
## [1] 12
```

```
memory.limit()
```

```
## [1] 65471
```

```
# Script settings  
knitr::opts_chunk$set(cache = TRUE)  
script_start <- Sys.time()  
  
# Load required packages  
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.2.1 --
```

```
## v ggplot2 3.1.1      v purrr   0.3.2  
## v tibble  2.1.1      v dplyr   0.8.0.1  
## v tidyr   0.8.3      v stringr 1.4.0  
## v readr   1.3.1      v forcats 0.4.0
```

```
## Warning: package 'ggplot2' was built under R version 3.5.3
```

```
## Warning: package 'tibble' was built under R version 3.5.3
```

```
## Warning: package 'tidyr' was built under R version 3.5.3
```

```
## Warning: package 'readr' was built under R version 3.5.3
```

```
## Warning: package 'purrr' was built under R version 3.5.3
```

```
## Warning: package 'dplyr' was built under R version 3.5.3
```

```
## Warning: package 'stringr' was built under R version 3.5.3
```

```
## Warning: package 'forcats' was built under R version 3.5.3
```

```
## -- Conflicts ----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()     masks stats::lag()
```

```
library(lubridate)
```

```
##  
## Attaching package: 'lubridate'  
  
## The following object is masked from 'package:base':  
##  
##     date
```

```
library(parallel)  
library(doParallel)
```

```
## Warning: package 'doParallel' was built under R version 3.5.3
```

```
## Loading required package: foreach
```

```
##  
## Attaching package: 'foreach'
```

```
## The following objects are masked from 'package:purrr':  
##  
##     accumulate, when
```

```
## Loading required package: iterators
```

```
## Warning: package 'iterators' was built under R version 3.5.1
```

Loading packages and data

The dataset was downloaded from the GroupLens website using the course provided script and stored locally. Information about the data was found on the [MovieLens README] (<http://files.grouplens.org/datasets/movielens/ml-10m-README.html>)

```
# Load locally stored data  
edx_original <- readRDS("edx.RData")  
print(str(edx_original))
```

```
## 'data.frame':    9000055 obs. of  6 variables:  
## $ userId   : int  1 1 1 1 1 1 1 1 1 1 ...  
## $ movieId  : num  122 185 292 316 329 355 356 362 364 370 ...  
## $ rating   : num  5 5 5 5 5 5 5 5 5 5 ...  
## $ timestamp: int  838985046 838983525 838983421 838983392 838983392 838984474 838983653 838984885 8...  
## $ title    : chr   "Boomerang (1992)" "Net, The (1995)" "Outbreak (1995)" "Stargate (1994)" ...  
## $ genres   : chr   "Comedy|Romance" "Action|Crime|Thriller" "Action|Drama|Sci-Fi|Thriller" "Action|A...  
## NULL
```

```
validation_original <- readRDS("validation.RData")
print(str(validation_original))
```

```
## 'data.frame':    999999 obs. of  6 variables:
## $ userId      : int  1 1 1 2 2 2 3 3 4 4 ...
## $ movieId     : num  231 480 586 151 858 ...
## $ rating      : num  5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ timestamp: int  838983392 838983653 838984068 868246450 868245645 868245920 1136075494 1133571200
## $ title       : chr   "Dumb & Dumber (1994)" "Jurassic Park (1993)" "Home Alone (1990)" "Rob Roy (1995)"
## $ genres      : chr   "Comedy" "Action|Adventure|Sci-Fi|Thriller" "Children|Comedy" "Action|Drama|Roman
## NULL
```

```
# Make sure movies and users are in both train and test sets ### USERS DIFFER
train_data <- edx_original
print(str(train_data))
```

```
## 'data.frame':    9000055 obs. of  6 variables:
## $ userId      : int  1 1 1 1 1 1 1 1 1 1 ...
## $ movieId     : num  122 185 292 316 329 355 356 362 364 370 ...
## $ rating      : num  5 5 5 5 5 5 5 5 5 5 ...
## $ timestamp: int  838985046 838983525 838983421 838983392 838983392 838984474 838983653 838984885 8
## $ title       : chr   "Boomerang (1992)" "Net, The (1995)" "Outbreak (1995)" "Stargate (1994)" ...
## $ genres      : chr   "Comedy|Romance" "Action|Crime|Thriller" "Action|Drama|Sci-Fi|Thriller" "Action|A
## NULL
```

```
test_data <- validation_original %>%
  semi_join(train_data, by = "movieId") %>%
  semi_join(train_data, by = "userId")
print(str(test_data))
```

```
## 'data.frame':    999999 obs. of  6 variables:
## $ userId      : int  1 1 1 2 2 2 3 3 4 4 ...
## $ movieId     : num  231 480 586 151 858 ...
## $ rating      : num  5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ timestamp: int  838983392 838983653 838984068 868246450 868245645 868245920 1136075494 1133571200
## $ title       : chr   "Dumb & Dumber (1994)" "Jurassic Park (1993)" "Home Alone (1990)" "Rob Roy (1995)"
## $ genres      : chr   "Comedy" "Action|Adventure|Sci-Fi|Thriller" "Children|Comedy" "Action|Drama|Roman
## NULL
```

```
all_data <- setNames(list(train_data, test_data), c("train_data", "test_data"))
```

Preparing data

Calculating statistics

```
## User statistics
calc_user_stats <- function(rating_data){
```

```

rating_data %>% group_by(userId) %>%
  summarise(user_ratings = n(),
            user_median = median(rating),
            user_mean = mean(rating),
            user_sd = sd(rating)) %>%
  select(userId, user_median, user_mean, user_sd) %>%
  unique()
}

user_stats <- parLapply(pcl, rating_data, calc_user_stats)

# Movie rating statistics
calc_movie_stats <- function(rating_data){

  rating_data %>% group_by(movieId) %>%
    summarise(movie_ratings = n(),
              movie_median = median(rating),
              movie_mean = mean(rating),
              movie_sd = sd(rating)) %>%
    select(movieId, movie_median, movie_mean, movie_sd) %>%
    unique()
}

movie_stats <- parLapply(pcl, rating_data, calc_movie_stats)

# Merge statistics back
merge_data <- list(list(rating_data[[1]], user_stats[[1]], movie_stats[[1]]), # train data
                  list(rating_data[[2]], user_stats[[2]], movie_stats[[2]])) # test data

merge_stats <- function(data_stats){

  rating_data <- data_stats[[1]]
  user_stats <- data_stats[[2]]
  movie_stats <- data_stats[[3]]

  rating_data %>% inner_join(user_stats, by = "userId") %>%
    inner_join(movie_stats, by = "movieId")
}

rating_data <- parLapply(pcl, merge_data, merge_stats)

## Rating characteristics
calc_rating_stats <- function(rating_data){
  rating_mean <- mean(rating_data$rating)

  rating_data <- rating_data %>% mutate(rating_mean_diff = rating - rating_mean) %>%
    mutate(movie_mean_diff = rating - movie_mean) %>%
    mutate(user_mean_diff = rating - user_mean)
}

```

```
rating_data <- parLapply(pcl, rating_data, calc_rating_stats)
stopCluster(pcl)
```

Derived models

```
train_ratings <- rating_data[[1]]
test_ratings <- rating_data[[2]]

saveRDS(train_ratings, "train_ratings.rds")
saveRDS(test_ratings, "test_ratings.rds")
```

```
# RMSE loss function to evaluate models
RMSE <- function(predicted_ratings, true_ratings){
  sqrt(mean((true_ratings - predicted_ratings)^2))
}
```

```
# Derived models
```

```
# Naive model using average rating
train_mean <- mean(train_ratings$rating)
train_mean
```

```
## [1] 3.512465
```

```
naive_rmse <- RMSE(test_ratings$rating, train_mean)
```

```
derived_results <- tibble(method = "Average training rating", RMSE = naive_rmse)
```

```
# Model using fixed number
fixed_number <- rep(2.5, nrow(test_ratings))
```

```
fixed_rmse <- RMSE(test_ratings$rating, fixed_number)
```

```
derived_results <- bind_rows(derived_results,
  tibble(method = "Fixed number 2.5",
    RMSE = fixed_rmse ))
```

```
# Derived predictions
```

```
movie_pred <- test_ratings %>% group_by(movieId) %>%
  # Rating effect (effect of rating scale used; looking at differences from rating mean)
  summarise(rating_mean_diff_avg = train_mean + mean(rating_mean_diff),
  # Movie effect (effect of general movie popularity; looking at differences between user
  movie_mean_diff_avg = train_mean + mean(movie_mean_diff),
  # User effect (effect of user rating behaviour; looking at differences between rating and
  user_mean_diff_avg = train_mean + mean(user_mean_diff),
  # User-movie effect (effect of user rating behaviour; looking at differences between user
  user_movie_mean_diff_avg = train_mean + mean(mean(movie_mean_d
```

```

der_pred <- test_ratings %>% inner_join(movie_pred, by = "movieId")

# Calculate RMSEs
rating_effect_rmse <- RMSE(test_ratings$rating,
                           der_pred$rating_mean_diff_avg)
movie_mean_rmse <- RMSE(test_ratings$rating,
                        test_ratings$movie_mean)
movie_effect_rmse <- RMSE(test_ratings$rating,
                          der_pred$movie_mean_diff_avg)
user_effect_rmse <- RMSE(test_ratings$rating,
                         der_pred$user_mean_diff_avg)
user_movie_effect_rmse <- RMSE(test_ratings$rating,
                               der_pred$user_movie_mean_diff_avg)

# Print RMSEs

derived_results <- bind_rows(derived_results,
                             tibble(method = "Rating effect",
                                     RMSE = rating_effect_rmse))
derived_results <- bind_rows(derived_results,
                             tibble(method = "Movie mean",
                                     RMSE = movie_mean_rmse))
derived_results <- bind_rows(derived_results,
                             tibble(method = "Movie effect",
                                     RMSE = movie_effect_rmse))
derived_results <- bind_rows(derived_results,
                             tibble(method = "User effect",
                                     RMSE = user_effect_rmse))
derived_results <- bind_rows(derived_results,
                             tibble(method = "User-movie effect",
                                     RMSE = user_movie_effect_rmse))

saveRDS(derived_results, "derived_results.rds")
print(derived_results %>% arrange(RMSE))

```

```

## # A tibble: 7 x 2
##   method          RMSE
##   <chr>          <dbl>
## 1 Movie mean      0.938
## 2 Rating effect   0.938
## 3 User effect     0.945
## 4 Average training rating 1.06
## 5 Movie effect    1.06
## 6 User-movie effect 1.06
## 7 Fixed number 2.5  1.47

```

Train linear models

```

# Load locally stored data
train_ratings <- readRDS("train_ratings.rds")

```

```
test_ratings <- readRDS("test_ratings.rds")
```

```
# Free up memory
```

```
rm(edx_original,
   validation_original,
   train_data,
   test_data,
   all_data,
   user_data,
   movie_data,
   rating_data,
   merge_data,
   der_pred)
gc() # garbage collection
```

```
##           used (Mb) gc trigger (Mb) max used (Mb)
## Ncells  1102650 58.9  11587712 618.9 14484641 773.6
## Vcells 127185803 970.4 389220886 2969.6 483144684 3686.2
```

```
# Set up parallel processing
```

```
no_cores <- detectCores() - 2
tcl <- makePSOCKcluster(no_cores)
registerDoParallel(tcl)
invisible(clusterEvalQ(tcl, library(caret)))

print(paste("Memory size before garbage collection: ", memory.size()))
```

```
## [1] "Memory size before garbage collection: 1569.26"
```

```
gc()
```

```
##           used (Mb) gc trigger (Mb) max used (Mb)
## Ncells  1105161 59.1   9270169 495.1 14484641 773.6
## Vcells 127191574 970.4 389220886 2969.6 483144684 3686.2
```

```
print(paste("Memory size after garbage collection: ", memory.size()))
```

```
## [1] "Memory size after garbage collection: 1062.05"
```

```
## Define models
```

```
user_formulas <- c("rating ~ user_median",
                  "rating ~ user_median + user_mean",
                  "rating ~ user_median + user_mean + user_sd")

movie_formulas <- c("rating ~ movie_median",
                  "rating ~ movie_median + movie_mean",
                  "rating ~ movie_median + movie_mean + movie_sd")

combined_formulas <- c("rating ~ movie_mean + user_mean",
                     "rating ~ movie_median + user_median")
```

```

full_formulas <- c("rating ~ movie_median + movie_mean + user_median + user_mean",
                  "rating ~ movie_median + movie_mean + movie_sd + user_median + user_mean + user_sd")

## Model training

e <- simpleError("Catch error")

train_lm <- function(lm_formula, lm_data){

  print(paste("Model formula: ", lm_formula))

  print(system.time({
    model <- tryCatch({
      train(as.formula(lm_formula),
            data = lm_data, method = "lm",
            na.action = na.omit)
    }, error = function(e) e, finally = print("Made it!"))
  }))

  # Full train objects are GBs large; save only the final model coefficients
  model_coef <- model$finalModel$coefficients

  # How's our memory?
  print(paste("Memory size before garbage collection: ", memory.size()))
  gc()
  print(paste("Memory size after garbage collection: ", memory.size()))

  # Save model locally
  saveRDS(model_coef, paste0("models/", lm_formula, ".rds"))
  # Return results
  print(model_coef)
  model
}

# Train and save locally
user_fits <- parLapply(tcl, user_formulas, train_lm, lm_data = train_ratings)
#saveRDS(user_fits, "user_fits.rds")
rm(user_fits)

movie_fits <- parLapply(tcl, movie_formulas, train_lm, lm_data = train_ratings)
#saveRDS(movie_fits, "movie_fits.rds")
rm(movie_fits)

combined_fits <- parLapply(tcl, combined_formulas, train_lm, lm_data = train_ratings)
#saveRDS(combined_fits, "combined_fits.rds")
rm(combined_fits)

full_fits <- parLapply(tcl, full_formulas, train_lm, lm_data = train_ratings)
#saveRDS(full_fits, "full_fits.rds")
rm(full_fits)

# Stop parallel processing
stopImplicitCluster()

```



```
# Clean up before next step
gc()
```

```
##          used   (Mb) gc trigger   (Mb)    max used   (Mb)
## Ncells  2272729 121.4  28297752 1511.3   20249735 1081.5
## Vcells 133315388 1017.2 5560351332 42422.2 8437143122 64370.3
```

Final model predictions and results

```
# Read model coeffecients and calculate RMSE
```

```
calculate_results <- function(model_formulas, test_ratings){
  str(test_ratings)
```

```
  for(m in 1:length(model_formulas)){
```

```
    #print("Current model: ")
    #print(model_formulas[[m]])
```

```
    coefs <- readRDS(paste0("models/", model_formulas[[m]], ".rds"))
    #print("Coefficients read: ")
    #print(coefs)
```

```
    prediction <- rep(coefs[[1]], nrow(test_ratings)) # start with the intercept
    #print("Prediction (intercept): ")
    #str(prediction)
```

```
    for(c in 2:length(coefs)){ # then go through each coefficient
      #print("Current coefficient: ")
      #print(coefs[c])
```

```
      ce <- rep(coefs[[c]], nrow(test_ratings))
      #str(ce)
```

```
      #print(paste("Column name: ", names(coefs[c])))
      #str(test_ratings[, names(coefs[c])))
```

```
      prediction <- prediction + (ce * test_ratings[, names(coefs[c])]) # and add coef * value to predi
```

```
    }
```

```
    #print("Final prediction: ")
    #str(prediction)
    #print("Test ratings: ")
    #str(test_ratings$rating)
```

```
    RMSE <- function(predicted_ratings, true_ratings){
      #print("Calculating RMSE!")
      #str(predicted_ratings)
      #print(predicted_ratings[1:10])
```

```

    #str(true_ratings)
    #print(true_ratings[1:10])

    rmse <- sqrt(mean((true_ratings - predicted_ratings)^2, na.rm = TRUE))

    #print("RMSE!")
    #str(rmse)
    #print(rmse)

    rmse
  }

  model_rmse <- RMSE(prediction, test_ratings$rating)
  #print(paste("Model RMSE: ", model_rmse))

  #print(paste(model_formulas[[m]], model_rmse))

  results <- data.frame(model = model_formulas[m], rmse = model_rmse)
}

results
}

user_results <- lapply(user_formulas, calculate_results, test_ratings = test_ratings)

```

```

## 'data.frame': 999999 obs. of 13 variables:
## $ userId : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating : num 5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year : num 1996 1996 1996 1997 1997 ...
## $ user_median : num 5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean : num 5 5 5 2.67 2.67 ...
## $ user_sd : num 0 0 0 0.577 0.577 ...
## $ movie_median : num 3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean : num 2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd : num 1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num 1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num 2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff : num 0 0 0 0.333 -0.667 ...
## 'data.frame': 999999 obs. of 13 variables:
## $ userId : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating : num 5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year : num 1996 1996 1996 1997 1997 ...
## $ user_median : num 5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean : num 5 5 5 2.67 2.67 ...
## $ user_sd : num 0 0 0 0.577 0.577 ...
## $ movie_median : num 3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean : num 2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd : num 1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num 1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num 2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff : num 0 0 0 0.333 -0.667 ...

```

```
## 'data.frame': 999999 obs. of 13 variables:
## $ userId : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating : num 5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year : num 1996 1996 1996 1997 1997 ...
## $ user_median : num 5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean : num 5 5 5 2.67 2.67 ...
## $ user_sd : num 0 0 0 0.577 0.577 ...
## $ movie_median : num 3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean : num 2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd : num 1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num 1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num 2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff : num 0 0 0 0.333 -0.667 ...
```

```
movie_results <- lapply(movie_formulas, calculate_results, test_ratings = test_ratings)
```

```
## 'data.frame': 999999 obs. of 13 variables:
## $ userId : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating : num 5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year : num 1996 1996 1996 1997 1997 ...
## $ user_median : num 5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean : num 5 5 5 2.67 2.67 ...
## $ user_sd : num 0 0 0 0.577 0.577 ...
## $ movie_median : num 3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean : num 2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd : num 1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num 1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num 2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff : num 0 0 0 0.333 -0.667 ...
## 'data.frame': 999999 obs. of 13 variables:
## $ userId : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating : num 5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year : num 1996 1996 1996 1997 1997 ...
## $ user_median : num 5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean : num 5 5 5 2.67 2.67 ...
## $ user_sd : num 0 0 0 0.577 0.577 ...
## $ movie_median : num 3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean : num 2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd : num 1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num 1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num 2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff : num 0 0 0 0.333 -0.667 ...
## 'data.frame': 999999 obs. of 13 variables:
## $ userId : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating : num 5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year : num 1996 1996 1996 1997 1997 ...
## $ user_median : num 5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean : num 5 5 5 2.67 2.67 ...
## $ user_sd : num 0 0 0 0.577 0.577 ...
## $ movie_median : num 3 4 3 4 5 3 4 4 4 3 ...
```

```
## $ movie_mean      : num  2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd        : num  1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num  1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num  2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff  : num  0 0 0 0.333 -0.667 ...
```

```
combined_results <- lapply(combined_formulas, calculate_results, test_ratings = test_ratings)
```

```
## 'data.frame': 999999 obs. of 13 variables:
## $ userId          : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId         : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating          : num  5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year     : num  1996 1996 1996 1997 1997 ...
## $ user_median     : num  5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean       : num  5 5 5 2.67 2.67 ...
## $ user_sd         : num  0 0 0 0.577 0.577 ...
## $ movie_median    : num  3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean      : num  2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd        : num  1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num  1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num  2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff  : num  0 0 0 0.333 -0.667 ...
## 'data.frame': 999999 obs. of 13 variables:
## $ userId          : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId         : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating          : num  5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year     : num  1996 1996 1996 1997 1997 ...
## $ user_median     : num  5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean       : num  5 5 5 2.67 2.67 ...
## $ user_sd         : num  0 0 0 0.577 0.577 ...
## $ movie_median    : num  3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean      : num  2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd        : num  1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num  1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num  2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff  : num  0 0 0 0.333 -0.667 ...
```

```
full_results <- lapply(full_formulas, calculate_results, test_ratings = test_ratings)
```

```
## 'data.frame': 999999 obs. of 13 variables:
## $ userId          : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId         : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 ...
## $ rating          : num  5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year     : num  1996 1996 1996 1997 1997 ...
## $ user_median     : num  5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean       : num  5 5 5 2.67 2.67 ...
## $ user_sd         : num  0 0 0 0.577 0.577 ...
## $ movie_median    : num  3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean      : num  2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd        : num  1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num  1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num  2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff  : num  0 0 0 0.333 -0.667 ...
```

```
## 'data.frame': 999999 obs. of 13 variables:
## $ userId : Factor w/ 68534 levels "1","2","3","4",...: 1 1 1 2 2 2 3 3 4 4 ...
## $ movieId : Factor w/ 9809 levels "1","2","3","4",...: 228 475 579 149 834 1477 583 4777 34 4 ...
## $ rating : num 5 5 5 3 2 3 3.5 4.5 5 3 ...
## $ rating_year : num 1996 1996 1996 1997 1997 ...
## $ user_median : num 5 5 5 3 3 3 4 4 3 3 ...
## $ user_mean : num 5 5 5 2.67 2.67 ...
## $ user_sd : num 0 0 0 0.577 0.577 ...
## $ movie_median : num 3 4 3 4 5 3 4 4 4 3 ...
## $ movie_mean : num 2.95 3.64 3.07 3.57 4.41 ...
## $ movie_sd : num 1.222 0.943 0.98 0.913 0.795 ...
## $ rating_mean_diff: num 1.488 1.488 1.488 -0.512 -1.512 ...
## $ movie_mean_diff : num 2.047 1.356 1.925 -0.572 -2.413 ...
## $ user_mean_diff : num 0 0 0 0.333 -0.667 ...
```

```
# Save results locally and print
saveRDS(user_results, "user_results.rds")
saveRDS(movie_results, "movie_results.rds")
saveRDS(combined_results, "combined_results.rds")
saveRDS(full_results, "full_results.rds")

# Lets have a look
lm_results <- bind_rows(c(user_results, movie_results, combined_results, full_results))
```

```
## Warning in bind_rows(x, .id): Unequal factor levels: coercing to character
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,
## coercing into character vector
```

```
## Warning in bind_rows(x, .id): binding character and factor vector,  
## coercing into character vector
```

```
lm_results %>% arrange(rmse)
```

```
##                                     model  
## 1          rating ~ movie_median + movie_mean + user_median + user_mean  
## 2                                     rating ~ movie_mean + user_mean  
## 3 rating ~ movie_median + movie_mean + movie_sd + user_median + user_mean + user_sd  
## 4          rating ~ movie_median + user_median  
## 5          rating ~ movie_median + movie_mean  
## 6          rating ~ movie_median + movie_mean + movie_sd  
## 7          rating ~ user_median + user_mean  
## 8          rating ~ user_median + user_mean + user_sd  
## 9          rating ~ movie_median  
## 10         rating ~ user_median  
##      rmse  
## 1  0.8452112  
## 2  0.8452248  
## 3  0.8467370  
## 4  0.8843992  
## 5  0.9383091  
## 6  0.9386363  
## 7  0.9395215  
## 8  0.9413628  
## 9  0.9578211  
## 10 0.9679871
```

```
saveRDS(lm_results, "lm_results.rds")
```

```
# Clean up before next step  
gc()
```

```
##          used   (Mb) gc trigger   (Mb)    max used   (Mb)  
## Ncells  2276497 121.6  22638201 1209.1  20249735 1081.5  
## Vcells 133328768 1017.3 4448281065 33937.7 8437143122 64370.3
```

Automatic Machine Learning with H2O.ai

```
# Shall we explore auto ML?
```

```
library(h2o)
```

```
h2o.init()
```

```
##  
## H2O is not running yet, starting it now...  
##  
## Note: In case of errors look at the following log files:
```

```
## C:\Users\Codrin\AppData\Local\Temp\RtmpWqEYjc/h2o_Codrin_started_from_r.out
## C:\Users\Codrin\AppData\Local\Temp\RtmpWqEYjc/h2o_Codrin_started_from_r.err
##
##
## Starting H2O JVM and connecting: Connection successful!
##
## R is connected to the H2O cluster:
## H2O cluster uptime: 1 seconds 437 milliseconds
## H2O cluster timezone: Europe/Berlin
## H2O data parsing timezone: UTC
## H2O cluster version: 3.24.0.3
## H2O cluster version age: 23 days
## H2O cluster name: H2O_started_from_R_Codrin_hvm578
## H2O cluster total nodes: 1
## H2O cluster total memory: 14.21 GB
## H2O cluster total cores: 12
## H2O cluster allowed cores: 12
## H2O cluster healthy: TRUE
## H2O Connection ip: localhost
## H2O Connection port: 54321
## H2O Connection proxy: NA
## H2O Internal Security: FALSE
## H2O API Extensions: Amazon S3, Algos, AutoML, Core V3, Core V4
## R Version: R version 3.5.0 (2018-04-23)
```

```
h2o.no_progress()
```

```
# Import a sample binary outcome train/test set into H2O
train <- as.h2o(readRDS("edx.RData"))
test <- as.h2o(readRDS("validation.RData"))

# Identify predictors and response
y <- "rating"
x <- setdiff(names(train), y)

# Run AutoML for 10 base models (limited to 1 hour max runtime by default)
aml <- h2o.automl(x = x, y = y,
  training_frame = train,
  validation_frame = test,
  max_models = 10,
  seed = 1,
  stopping_metric = "RMSE",
  sort_metric = "RMSE")

# View the AutoML Leaderboard
lb <- aml@leaderboard
autoML_results <- as.data.frame(lb) %>% select(model_id, rmse)

saveRDS(autoML_results, "autoML_results.rds")
print(autoML_results)
```

```
##                               model_id      rmse
## 1   StackedEnsemble_AllModels_AutoML_20190530_190225 0.9637047
## 2   StackedEnsemble_BestOfFamily_AutoML_20190530_190225 0.9638613
```

```
## 3          GBM_5_AutoML_20190530_190225 0.9763070
## 4          DRF_1_AutoML_20190530_190225 0.9889825
## 5          XRT_1_AutoML_20190530_190225 0.9930478
## 6          GBM_4_AutoML_20190530_190225 1.0087012
## 7          GBM_3_AutoML_20190530_190225 1.0196827
## 8          GBM_2_AutoML_20190530_190225 1.0240631
## 9          GBM_1_AutoML_20190530_190225 1.0277414
## 10         DeepLearning_1_AutoML_20190530_190225 1.0461689
## 11         GBM_grid_1_AutoML_20190530_190225_model_1 1.0570559
## 12         GLM_grid_1_AutoML_20190530_190225_model_1 1.0596654
```

```
# How long did the whole script take?
```

```
script_end <- Sys.time()
```

```
print(paste("Total script running time: ", round(difftime(script_end, script_start, units = "mins"), 1))
```

```
## [1] "Total script running time: 128.6 minutes"
```